

REGULARIZATION STRATEGIES IN DEEP LEARNING

CSE 676 – Project 2

-Sandesh Kumar Srivastava



University at Buffalo

Department of Computer Science
and Engineering


School of Engineering and Applied Sciences




Underfit vs Good Fit vs Overfit

- Underfit or High Bias **Solution: Make model more expressive/complex**
- Good fit or low bias and low variance
- Overfit or High Variance **Solution: Regularization**


John von Neumann famously said "With four parameters I can fit an elephant, and with five I can make him wiggle his trunk."



Underfitted




Good Fit/Robust



Overfitted

Overfitting in real life example



Regularization Strategies

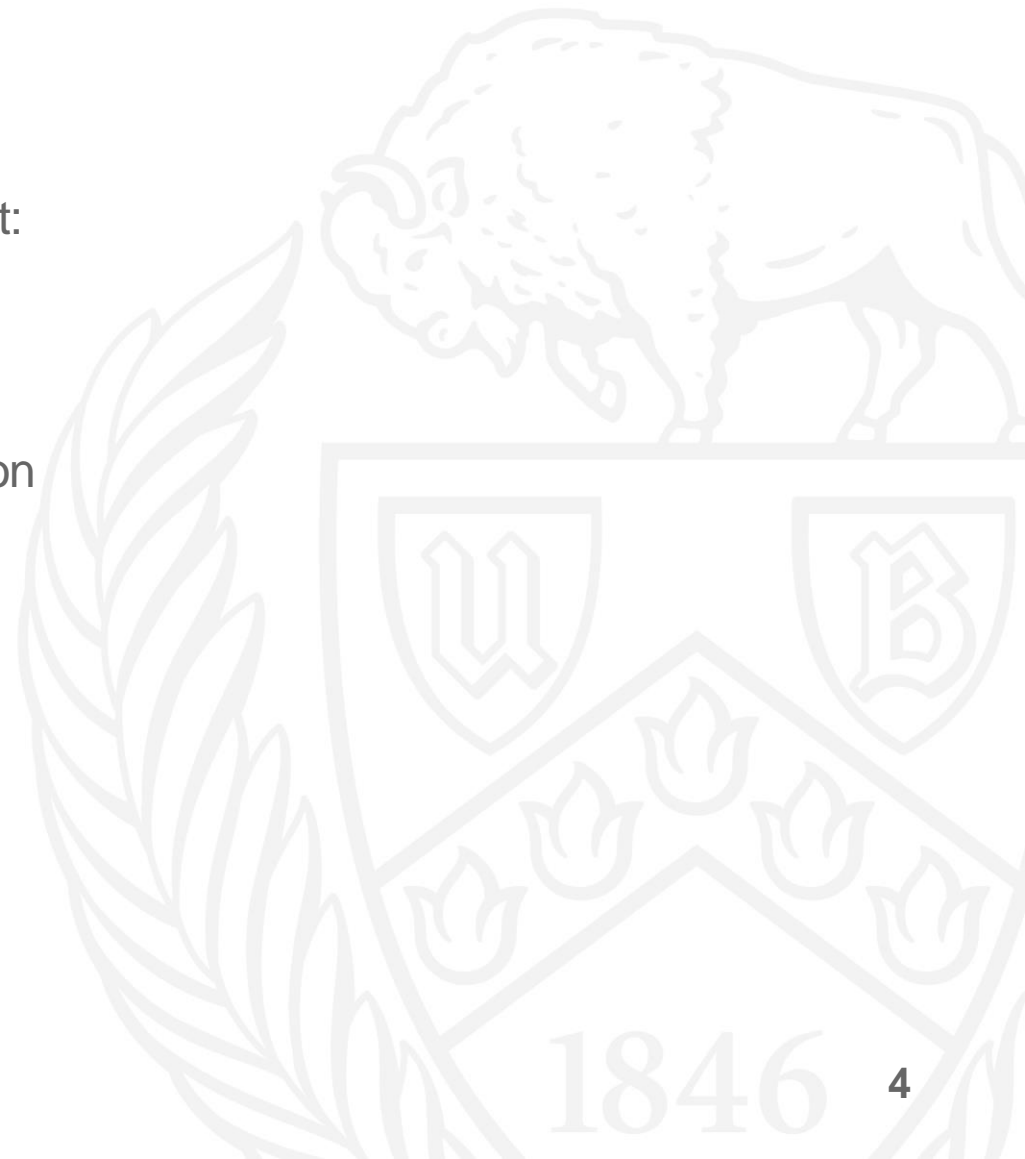
Methods to reduce generalization error(error on unseen data)

-> But not the training error

-> Even at the expense of training error

Different methods exist:

- L1 norm
- L2 norm
- Data Set Augmentation
- Noise Robustness
- Early Stopping
- Dropout
- Adversarial Training ..



Regularization Strategies

L1 norm:

$$\tilde{J}(\Theta; X, y) = J(\Theta; X, y) + \alpha \Omega(\Theta) \quad (1)$$

$$\Omega(\Theta) = \|w\|_1 = \sum |w_i|_1 \quad (2)$$

$$\tilde{J}(\Theta; X, y) = J(\Theta; X, y) + \alpha \sum_i |w_i|_1 \quad (3)$$

```
Conv2D(filters=f_5x5, kernel_size=(5,5),
padding='same', activation='relu', kernel
_regularizer=l1(l1=0.01), bias_regularize
r=l1(1e-4), activity_regularizer=l1(1e-
5))
```

L2 norm:

$$\tilde{J}(\Theta; X, y) = J(\Theta; X, y) + \alpha \Omega(\Theta) \quad (1)$$

$$\Omega(\Theta) = \frac{1}{2} \|w\|_2^2 \quad (4)$$

$$\tilde{J}(\Theta; X, y) = J(\Theta; X, y) + \alpha \frac{1}{2} \|w\|_2^2 \quad (5)$$

```
Conv2D(filters=f_3x3_r, kernel_size=(1,1)
, padding='same', activation='relu', kern
el_regularizer=l2(l2=0.001), bias_regular
izer=l2(1e-4),
activity_regularizer=l2(1e-5))
```


Regularization Strategies

Data Set Augmentation:

- Train the ML model on more data
- Transform the given input to obtain new input

```
datagen = ImageDataGenerator(  
    zca_epsilon=1e-06,  
    rotation_range=10,  
    width_shift_range=0.1,  
    height_shift_range=0.1,  
    horizontal_flip=True)
```

Noise Robustness:

- Noise can be applied at different levels to a ML model.
- If applied at input, it serves as a data augmentation.
- If applied to output layers, it helps to handle the mistakes made by ML model.

```
GaussianNoise(0.005)
```

Regularization Strategies

Early Stopping:

- Stop training process whenever there is not significant improvement on the validation data metrics

```
EarlyStopping(monitor='val_accuracy', patience=40)
```

Dropout:

- Technique similar to bagging.
- Randomly dropping some units by simply multiplying their output value to 0.

```
Dropout(0.5)
```

Adversarial Training:

- ML model trained on the generated adversarial examples.
- Fast Gradient Sign Method(FGSM)

$$x \rightarrow x + \epsilon \text{sign}(\nabla_x J(\Theta, x, y)) \quad (6)$$

```
gradient = tape.gradient(loss, image)
```

```
signed_grad = tf.sign.gradient)
```


```
adversarial = image + perturbations * epsilon
```

Implementation


- Using Keras library in Python.
- CIFAR-10 dataset(50k training and 10k test images of size 32x32x3).
- Inception_v2 like model.
- Adamax optimizer.
- ModelCheckPoint to store best weights.




Results - accuracy and loss plots




No Regularization




L1L2 norm




Early Stopping




L1 norm




Dataset Augmentation




Dropout



L2 norm





Noise Robustness



Adversarial Training (FGSM)

Results

| Regularization Strategy | Precision | Accuracy |
|-------------------------|-----------|----------|
| No regularization | 0.77896 | 0.7785 |
| L1 | 0.78399 | 0.7841 |
| L2 | 0.79902 | 0.8005 |
| L1L2 | 0.79272 | 0.7923 |
| Dataset Augmentation | 0.82933 | 0.8301 |
| Noise Robustness | 0.78409 | 0.7811 |
| Early Stopping | 0.77789 | 0.7703 |
| Dropout | 0.79751 | 0.7961 |
| Adversarial(FGSM) | 0.77787 | 0.777 |
| Combination | 0.84854 | 0.8481 |



Combination
(L2, Dataset Augmentation & Dropout)



← Achieved in 168 epochs compared to 200 for others

← Accuracy on adversarial test samples increased from 34% to 70%

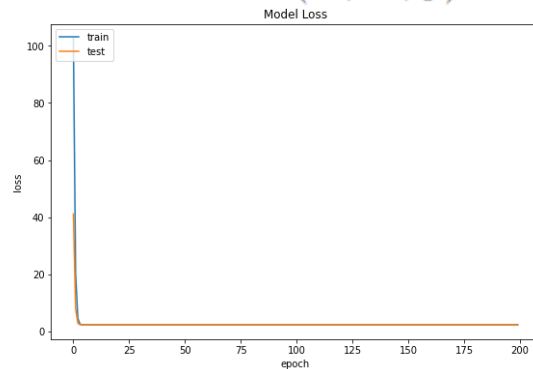

← Best performance

Results – finetuning regularization hyperparameters



| L2 kernel regularizer | Precision | Accuracy |
|-----------------------|-----------|----------|
| 1e-1 | 0.01 | 0.1 |
| 1e-2 | 0.504229 | 0.5173 |
| 1e-3 | 0.79902 | 0.8005 |
| 1e-4 | 0.77536 | 0.7759 |



$$\tilde{J}(\Theta; X, y) = J(\Theta; X, y) + \alpha\Omega(\Theta) \quad (1)$$



1e-1



1e-4

Conclusion

- Overfitting is a very common problem in deep neural networks.
- Regularization strategies help to solve the problem of overfitting.
- Selection of regularization strategy depends on the nature of problem being solved.
- Often multiple regularization strategies are combined to get best result.
- Finetuning of regularization parameters is required for optimal performance.

